

SYSC 4001 Assignment 3 (Part 1)

Faraz Aleboyeh 101311227, Rehma Muzammil 101268686

Part 1: https://github.com/rehmaaaa/SYSC4001_A3_P3

Part 2: https://github.com/farazaleboyeh/SYSC4001_A3_P2

SYSC 4001: Operating Systems, Fall 2025

1. Introduction

The purpose of Part 1 of the assignment was to build a small scheduler simulator. The system has a single CPU and a user space of 100 MB divided into 6 sizes: 40, 25, 15, 10, 8, and 2 MB.

We compared with three schedulers:

- 1. External Priorities (EP): non-preemptive, lower PID = higher priority.
- 2. Round Robin (RR): time-sharing with a fixed time quantum.
- 3. External Priorities + Round Robin (EP_RR): priority-based scheduling with preemption and a time quantum.

For each scheduler, we simulated 20 different traces and recorded the process state transitions in execution logs. From those logs we computed the following metrics:

- Throughput, Average Waiting Time, Average Turnaround Time, Average Response Time

2. Simulator Design: 2.1 Data Structures

We made fixed partitions in a table of structs: partition_number (1–6), size (40, 25, 15, 10, 8, 2 MB), code (string: "free", "init", or the program name / PID)

We also added a hook to log memory usage when a process transitions from NEW → READY.

Each process is represented by a PCB that stores the fields for this assignment:

- PID, arrival_time, processing_time (total CPU time), remaining_time (for preemption), io_freq and io_duration, state (NEW, READY, RUNNING, WAITING, TERMINATED), start_time and finish_time (for response and turnaround), partition_id

Queues are represented by std::vector<PCB>: ready_queue, wait_queue (for I/O), job_list

2.2 Process State Simulation

The main loop moves a discrete time variable also known as current_time which starts from 0 ms and at each step, the simulator:

- Brings new arrivals in (arrival_time == current_time)
 - Calls assign_memory(p), Changes state NEW → READY, pushes process into ready_queue and job_list, logs the transition in the execution file, logs memory usage in memory_log.txt
- If I/O completion time = current_time then WAITING → READY.
- Selects the next process if the CPU is idle.
- Reduces remaining_time, checks if it should go to I/O (RUNNING → WAITING), checks if it has finished (RUNNING → TERMINATED) and for RR and EP_RR, checks if the time quantum has expired (RUNNING → READY).

Each state change is recorded in an execution log: | time | pid | old_state | new_state | and these logs are used for calculations.

3. Scheduling Algorithms

3.1 External Priority (EP): EP is a non-preemptive priority scheduler. When the CPU is idle, the algorithm picks the READY process with the lowest PID and runs it until:

- It ends or It does an I/O and moves to WAITING.

No other process can preempt a running process even if it has a lower PID.

3.2 Round Robin (RR): RR uses a fixed time quantum, and the ready queue is like a FIFO:

- When a process is chosen, it goes for one time quantum but if it does not complete within the quantum and does not request I/O, it is preempted and is sent to the back of the queue but if it does I/O or finishes, it leaves the CPU.

3.3 External Priority + Round Robin (EP_RR)

- The highest priority (lowest PID) is chosen but the process still runs under a time quantum and if the quantum expires and the process isn't finished or is still waiting for an I/O then it is preempted and returns to READY.

This is used to keep the responsiveness of RR while still giving preference to higher-priority processes.

4. Setup and Metrics

We executed at least 20 simulation scenarios per scheduler:

- EP_traceX.txt, RR_traceX.txt, and EP_RR_traceX.txt contain the execution logs.
- For each log, we ran analyze_metrics.py to parse the transitions and compute.

- The results for every run are stored in metrics_summary.csv.

4.1 Metric Definitions

1. Throughput: Throughput= number of processes completed / total simulated time
2. Average Waiting Time: how long it stayed in READY and averaged over PIDs.
3. Average Turnaround Time: Turnaround=finish–arrival
4. Average Response Time (I/O): For I/O, we measured interval from RUNNING → WAITING to WAITING → READY, and then averaged it out

Overall,

- Throughput ranged from 0.024–0.064 processes per unit time for multi-process traces
- Average waiting times were around 10.67 ms or 15.33 ms, depending on the trace.
- Average turnaround times around 35–40 ms for shorter traces and around 64–69 ms for longer.
- Average response times for I/O were usually either about 2 ms or 3.45 ms.

Because the same trace sets were run under each scheduler algorithm, we can easily compare EP, RR, and EP_RR.

5. Results and Discussion

5.1 Throughput Comparison

For single-process traces all three schedulers got the same throughput:

- Because only one process exists, it runs to completion and has no competition.
- Throughput ≈ 0.1 , average wait time ≈ 0 , and turnaround time is the total CPU time.

For multi-process traces, throughput values go into 2 groups:

- 0.062–0.064 for shorter total run times
- 0.024–0.027 for longer simulations.

Across these EP, RR, and EP_RR all get similar throughput, which shows that the choice of scheduler does not change the total number of completed processes per unit time.

5.2 Waiting Time and Turnaround

- Average waiting times around 10.67 ms and turnaround ≈ 64.67 ms.
- RR and EP_RR runs have waiting ≈ 15.33 ms and turnaround ≈ 69.33 ms for the same number of processes.

EP reduces waiting time for at least some processes which lowers average turnaround a little bit. RR and EP_RR have more context switches which adds an extra delay before a process goes to the CPU, especially if there are many READY processes going through the quantum. However, the differences aren't huge because all three schedulers are in the same general range. Priority scheduling gives a small advantage in average wait/turnaround, but is balanced because RR and EP_RR are not horrible.

5.3 Response Time and I/O Behavior

The average response time is either 2.00 ms or 3.45 ms across all schedulers.

- I/O devices are always available (no queue)
- The main effect of the scheduler is when a process returns to READY after I/O and has to wait to be scheduled again.

Round Robin and EP_RR rotate through processes, so they can give more CPU slices to I/O-heavy processes, but this did not produce a large difference compared to EP. All three schedulers have low response times.

5.4 CPU-Bound vs I/O-Bound Processes

- CPU-bound processes benefit from priority scheduling (EP) because once they start running, they aren't preempted and can finish quickly which reduces their turnaround time which is why there were lower waiting times observed for EP traces.
- I/O-bound processes benefit from Round Robin behavior because if a process blocks for I/O, RR/EP_RR then it will make sure that other processes also get CPU slices, but when the I/O-bound process becomes READY again, it will be scheduled quickly into the rotation.

EP_RR preempts long-running processes using a quantum model that behaves similarly to pure RR which shows that all processes are similar in terms of CPU and I/O bursts.

6. Memory Usage (Bonus)

For the bonus part, we added logging of memory usage whenever a process is admitted (NEW → READY). So that means that the simulator assigns the process to one of the six fixed partitions according to its memory size and it logs a line to memory_log.txt in the format: TIME=<t> USED=<used_memory> FREE=<free_memory> USABLE=<usable_memory>

The output files memory_case_X.txt (one per trace) show the times at which processes start and how the memory layout changes. We found that large processes can only fit into the 40 MB or 25 MB partitions so if those are occupied then they wait even if smaller partitions are free which leads to internal fragmentation. For example, a 20 MB process in the 25 MB partition wastes 5 MB that no other process can use. This means that the number of admitted processes is bounded by the total partition sizes and their fit, which affects throughput and waiting time.

7. Conclusion

In this assignment, we built a scheduler simulator with:

- A single CPU and a fixed-partition memory model,
- Three schedulers: EP, RR, and EP_RR,
- Detailed logging of process state transitions and basic memory usage.

We ran at least 20 traces for each scheduler and computed throughput, average waiting time, turnaround time, and I/O response time using a Python analysis script.

The results show that:

- All three schedulers get a similar throughput.
- External Priorities (EP) reduces the average waiting and turnaround times for many traces by letting higher-priority processes run to completion without preemption.
- Round Robin (RR) and EP_RR gives more fairness with time-sharing, but affects waiting time due to the time quantum, but not a lot of performance loss.
- Response times for I/O are low and stable across algorithms, because I/O devices are always available and the fixed I/O duration always dominates.

Overall, the experiments prove that non-preemptive priority scheduling is slightly better for throughput and turnaround processes, Round Robin improves responsiveness, and a hybrid EP_RR scheduler balances both behaviors while preserving the process priorities.